

MSCIT 5210: Knowledge Discovery and Data Mining

Acknowledgement: Slides modified by Dr. Lei Chen based on the slides provided by Jiawei Han, Micheline Kamber, and Jian Pei

©2012 Han, Kamber & Pei. All rights reserved.

Chapter 9. Classification: Advanced Methods

- Classification by Backpropagation 
- Support Vector Machines
- Additional Topics Regarding Classification
- Summary


Biological Neural Systems

- Neuron switching time : $> 10^{-3}$ secs
- Number of neurons in the human brain: $\sim 10^{10}$
- Connections (synapses) per neuron : $\sim 10^4 - 10^5$
- Face recognition : 0.1 secs
- High degree of distributed and parallel computation
 - Highly fault tolerant
 - Highly efficient
 - Learning is key


Excerpt from Russell and Norvig

Brains

10^{11} neurons of > 20 types, 10^{14} synapses, 1ms–10ms cycle time
Signals are noisy “spike trains” of electrical potential



Modeling A Neuron on Computer




■ Computation:

- input signals \rightarrow input function(linear) \rightarrow activation function(nonlinear) \rightarrow output signal

Part 1. Perceptrons: Simple NN


inputs



x_i 's range: $[0, 1]$

$$y = \begin{cases} 1 & \text{if } a \geq \theta \\ 0 & \text{if } a < \theta \end{cases}$$

To be learned: W_i and θ



Converting θ To W_0

$$\sum_{i=1}^N W_i * X_i \geq \theta \quad (1)$$


$$\Leftrightarrow \sum_{i=1}^N W_i * X_i - \theta \geq 0 \quad (2)$$

$$\Leftrightarrow \sum_{i=1}^N W_i * X_i + (\theta) * (-1) \geq 0 \quad (3)$$


$$\Leftrightarrow \sum_{i=1}^N W_i * X_i + W_0 * X_0 \geq 0 \quad (4)$$

$$\Leftrightarrow \sum_{i=0}^N W_i * X_i \geq 0 \quad (5)$$

Threshold as Weight: W_0



Linear Separability



Logical AND


$$a = \sum_{i=0}^n w_i x_i$$

x_1	x_2	a	y
0	0	-1.5	0
0	1	-0.5	0
1	0	-0.5	0
1	1	0.5	1

t
0
0
0
1

$$y = \begin{cases} 1 & \text{if } a \geq 0 \\ 0 & \text{if } a < 0 \end{cases}$$

XOR cannot be separated!




Logical XOR

x_1	x_2	t	
0	0	0	
0	1	1	
1	0	1	
1	1	0	

Thus, one level neural network can only learn linear functions (straight lines)

Training the Perceptron

- Training set S of examples $\{\mathbf{x}, \mathbf{t}\}$
 - \mathbf{x} is an input vector and
 - \mathbf{T} the desired target vector (Teacher)
 - Example: Logical And 
- Iterative process
 - Present a training example x , compute network output y , compare output y with target t , adjust weights and thresholds
- Learning rule
 - Specifies how to change the weights W of the network as a function of the inputs x , output Y and target t .

x_1	x_2	t
0	0	0
0	1	0
1	0	0
1	1	1

Perceptron Learning Rule

$$w_i := w_i + \Delta w_i = w_i + \alpha (t-y) x_i \quad (i=1..n)$$

- The parameter α is called the *learning rate*.
 - In Han's book it is lower case L
 - It determines the magnitude of weight updates Δw_i .
- If the output is correct ($t=y$) the weights are not changed ($\Delta w_i = 0$).
- If the output is incorrect ($t \neq y$) the weights w_i are changed such that the output of the Perceptron for the new weights w'_i is *closer/further* to the input x_i .

Perceptron Training Algorithm

Repeat

for each training vector pair (\mathbf{x}, t)

evaluate the output y when \mathbf{x} is the input

if $y \neq t$ then

form a new weight vector \mathbf{w}' according
to $\mathbf{w}' = \mathbf{w} + \alpha (t - y) \mathbf{x}$

else α : set by the user; typically = 0.01


do nothing

end if

end for

Until fixed number of iterations; or error less than a
predefined value

Example: Learning the AND Function : Step 1.




x_1	x_2	t
0	0	0
0	1	0
1	0	0
1	1	1

W_0	W_1	W_2
0.5	0.5	0.5

$a = (-1) * 0.5 + 0 * 0.5 + 0 * 0.5 = -0.5$,
Thus, $y = 0$. Correct. No need to change W

$\alpha = 0.1$

Example: Learning the AND Function : Step 2.



x_1	x_2	t
0	0	0
0	1	0
1	0	0
1	1	1

$$\alpha = 0.1$$

W_0	W_1	W_2
0.5	0.5	0.5

$$a = (-1) * 0.5 + 0 * 0.5 + 1 * 0.5 = 0,$$

Thus, $y=1$. $t=0$, Wrong.

$$\Delta W_0 = 0.1 * (0 - 1) * (-1) = 0.1,$$

$$\Delta W_1 = 0.1 * (0 - 1) * (0) = 0$$

$$\Delta W_2 = 0.1 * (0 - 1) * (1) = -0.1$$

$$W_0 = 0.5 + 0.1 = 0.6$$

$$W_1 = 0.5 + 0 = 0.5$$

$$W_2 = 0.5 - 0.1 = 0.4$$

Example: Learning the AND Function : Step 3.

x_1	x_2	t
0	0	0
0	1	0
1	0	0
1	1	1



$\alpha = 0.1$

W_0	W_1	W_2
0.6	0.5	0.4

$a = (-1) * 0.6 + 1 * 0.5 + 0 * 0.4 = -0.1$,
Thus, $y = 0$. $t = 0$, Correct!

Example: Learning the AND Function : Step 2.

x_1	x_2	t
0	0	0
0	1	0
1	0	0
1	1	1




$\alpha = 0.1$

W_0	W_1	W_2
0.6	0.5	0.4

$a = (-1) * 0.6 + 1 * 0.5 + 1 * 0.4 = 0.3$,
Thus, $y = 1$. $t = 1$, Correct

Final Solution:



Logical AND

x_1	x_2	y
0	0	0
0	1	0
1	0	0
1	1	1

$$a = 0.5x_1 + 0.4x_2 - 0.6$$

$$y = \begin{cases} 1 & \text{if } a \geq 0 \\ 0 & \text{if } a < 0 \end{cases}$$



Perceptron Convergence Theorem

- The algorithm converges to the correct classification
 - if the training data is linearly separable
 - and learning rate is sufficiently small(Rosenblatt 1962).
- The final weights in the solution \mathbf{w} is not unique: there are many possible lines to separate the two classes.

Experiments


Perceptron learning contd.

Perceptron learning rule converges to a consistent function
for any linearly separable data set



Handwritten Recognition Example

Handwritten digit recognition

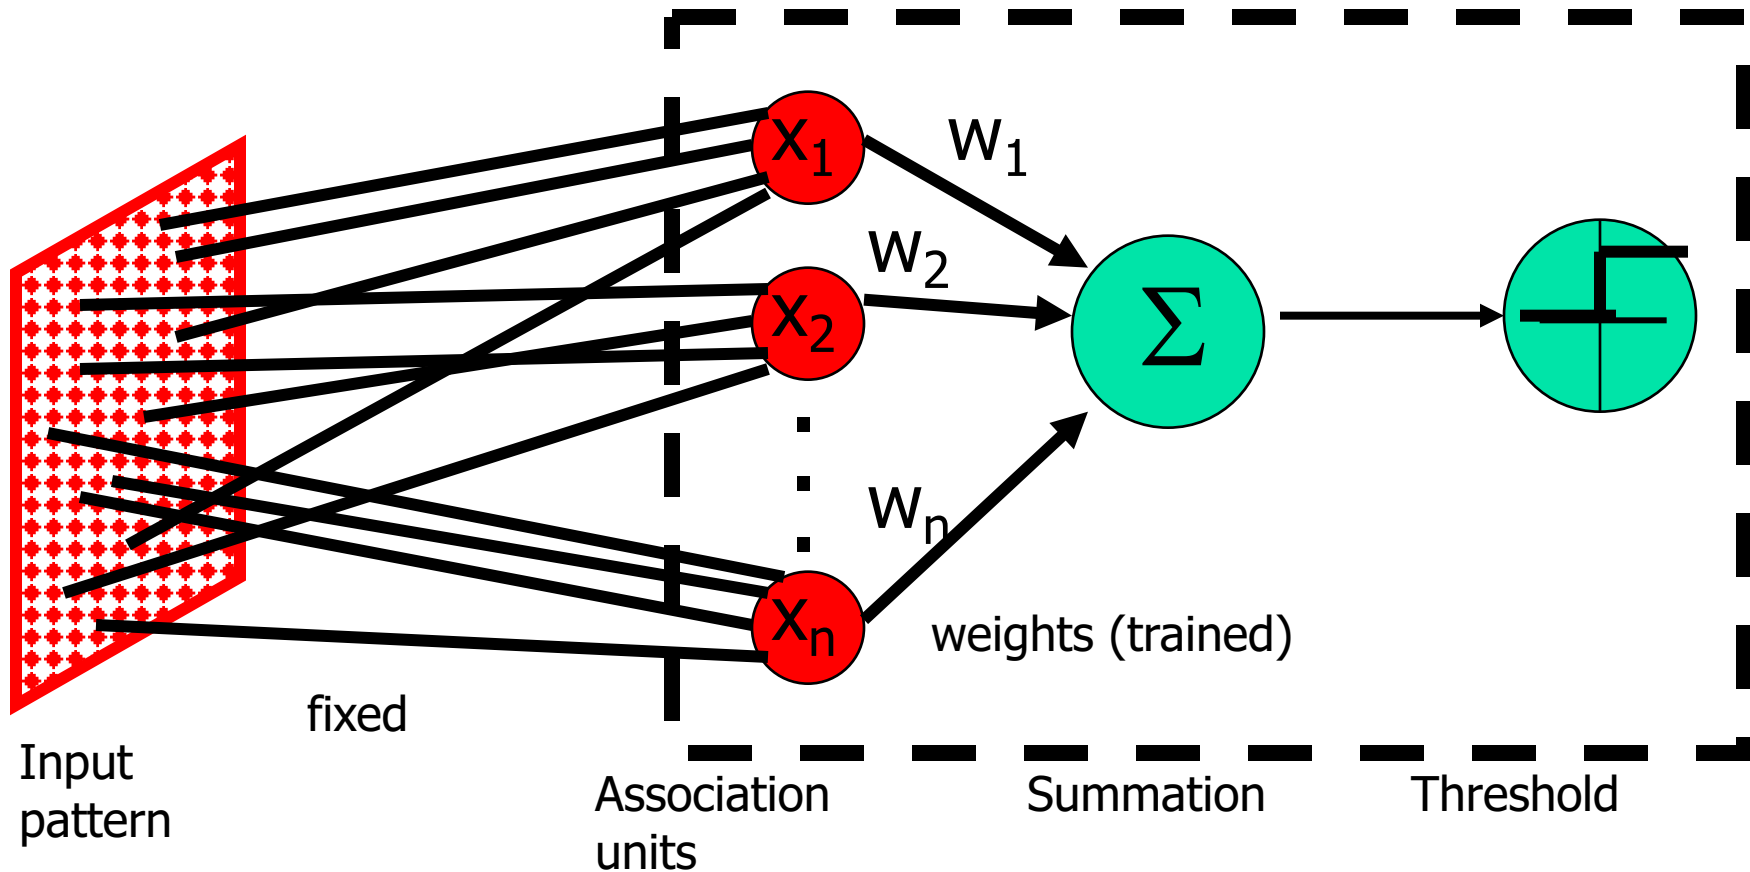


3-nearest-neighbor = 2.4% error

400-300-10 unit MLP = 1.6% error

LeNet: 768-192-30-10 unit MLP = 0.9%


Each letter \rightarrow one output unit y



Multiple Output Perceptrons


- Handwritten alphabetic character recognition
 - 26 classes : A,B,C...,Z
 - First output unit distinguishes between "A"s and "non-A"s, second output unit between "B"s and "non-B"s etc.

w_{ji} connects x_i with y_j




$$w'_{ji} = w_{ji} + \alpha (t_j - y_j) x_i$$

Part 2. Multi Layer Networks



Sigmoid-Function for Continuous Output

inputs




Output between 0 and 1 (when a = negative infinity, $O = 0$; when a = positive infinity, $O=1$).

Gradient Descent Learning Rule

- For each training example X ,
 - Let O be the output (between 0 and 1)
 - Let T be the correct target value
- Continuous output O
 - $a = w_1 x_1 + \dots + w_n x_n + \theta$
 - $O = 1/(1+e^{-a})$
- Train the w_i 's such that they minimize the squared error
 - $E[w_1, \dots, w_n] = \frac{1}{2} \sum_{k \in D} (T_k - O_k)^2$
where D is the set of training examples

Explanation: Gradient Descent Learning Rule



$$\Delta w_i = \alpha \frac{O_k (1 - O_k)}{x_i^k} (T_k - O_k)$$

learning rate $\rightarrow \alpha$

derivative of activation function $\rightarrow \frac{O_k (1 - O_k)}{x_i^k}$

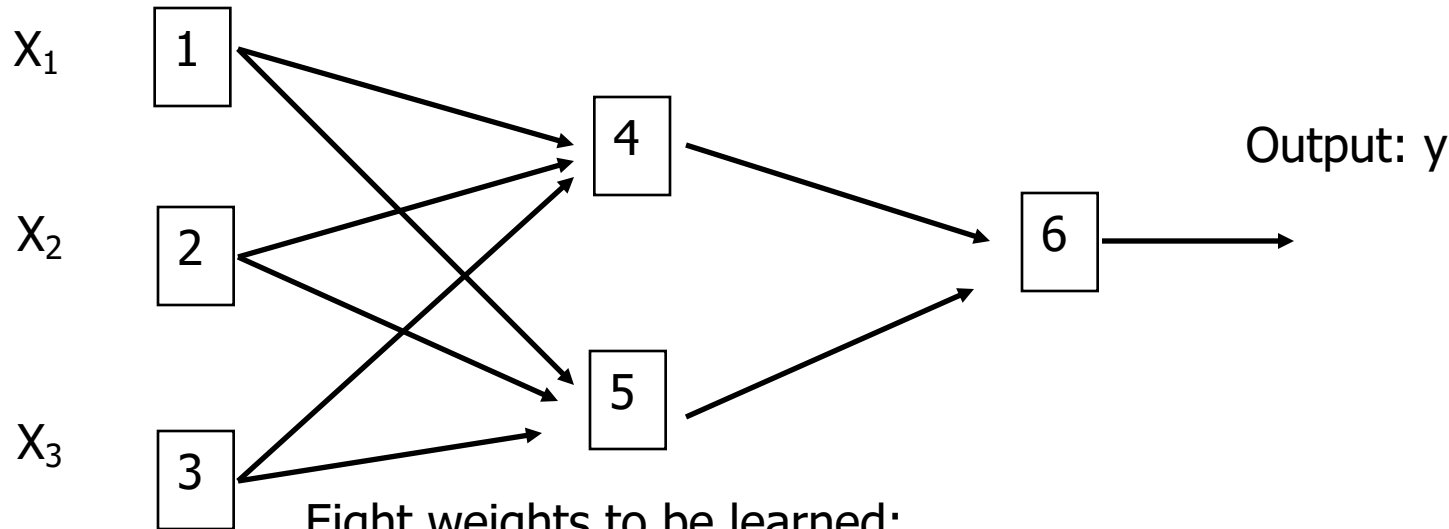
error δ_k of post-synaptic neuron $\rightarrow (T_k - O_k)$

activation of pre-synaptic neuron $\rightarrow x_i^k$

Backpropagation Algorithm (Hahn, Figure 9.5)

- Initialize each w_i to some small random value
 - Until the termination condition is met, Do
 - For each training example $\langle (x_1, \dots, x_n), t \rangle$ Do
 - Input the instance (x_1, \dots, x_n) to the network and compute the network outputs O_k
 - For each output unit k
 - $Err_k = O_k(1 - O_k)(t_k - O_k)$
 - For each hidden unit h
 - $Err_h = O_h(1 - O_h) \sum_k w_{h,k} Err_k$
 - For each network weight $w_{i,j}$ Do
 - $w_{i,j} = w_{i,j} + \Delta w_{i,j}$ where
 - $\Delta w_{i,j} = \alpha Err_j * O_i$
 - $\theta_j = \theta_j + \Delta \theta_j$ where
 - $\Delta \theta_j = \alpha Err_j$
- α : is learning rate, set by the user;

Example 6.9 (HK book, page 333)



Eight weights to be learned:

W_{ij} : W_{14} , W_{15} , ... W_{46} , W_{56} , ..., and θ

Training example:

x1	x2	x3	t
1	0	1	1

Learning rate: $=0.9$

Initial Weights: Randomly assigned (HK: Tables 7.3, 7.4)

x1	x2	x3	w14	w15	w24	w25	w34	w35	w46	w56	θ4	θ5	θ6
1	0	1	0.2	-0.3	0.4	0.1	-0.5	0.2	-0.3	-0.2	-0.4	0.2	0.1

Net input at unit 4:

$$X_1 * W_{14} + X_2 * W_{24} + X_3 * W_{34} + \theta_4 =$$
$$1 * (0.2) + 0 * 0.4 + 1 * (-0.5) + (-0.4) = -0.7$$

Output at unit 4:

$$\frac{1}{1 + e^{0.7}} = 0.332$$

Feed Forward: (Table 7.4)

- Continuing for units 5, 6 we get:
 - Output at unit 6 = 0.474

Calculating the error (Tables 7.5)

- Error at Unit 6: $(t-y)=(1-0.474)$

- Error to be backpropagated from unit 6:

$$y(1-y)(t-y) = (0.474)(1-0.474)(1-0.474) = 0.1311$$

- Weight update :

$$\Delta w_{46} = (l)Err_6 y_4 = 0.9 * (0.1311)(0.332)$$

$$w_{46} = w_{46} + (l)Err_6 y_4 = -0.3 + 0.9 * (0.1311)(0.332) \\ = -0.261$$

Weight update (Table 7.6)

Thus, new weights after training with $\{(1, 0, 1), t=1\}$:


w14	w15	w24	w25	w34	w35	w46	w56	θ_4	θ_5	θ_6
0.192	-0.306	0.4	0.1	-0.506	0.194	-0.261	-0.138	-0.408	0.194	0.218

- If there are more training examples, the same procedure is followed as above.
- Repeat the rest of the procedures.

Classification by Backpropagation

- Backpropagation: A **neural network** learning algorithm
- Started by psychologists and neurobiologists to develop and test computational analogues of neurons
- A neural network: A set of connected input/output units where each connection has a **weight** associated with it
- During the learning phase, the **network learns by adjusting the weights** so as to be able to predict the correct class label of the input tuples
- Also referred to as **connectionist learning** due to the connections between units

Neuron: A Hidden/Output Layer Unit



- An n -dimensional input vector \mathbf{x} is mapped into variable y by means of the scalar product and a nonlinear function mapping
- The inputs to unit are outputs from the previous layer. They are multiplied by their corresponding weights to form a weighted sum, which is added to the bias associated with unit. Then a nonlinear activation function is applied to it.


How A Multi-Layer Neural Network Works

- The **inputs** to the network correspond to the attributes measured for each training tuple
- Inputs are fed simultaneously into the units making up the **input layer**
- They are then weighted and fed simultaneously to a **hidden layer**
- The number of hidden layers is arbitrary, although usually only one
- The weighted outputs of the last hidden layer are input to units making up the **output layer**, which emits the network's prediction
- The network is **feed-forward**: None of the weights cycles back to an input unit or to an output unit of a previous layer
- From a statistical point of view, networks perform **nonlinear regression**: Given enough hidden units and enough training samples, they can closely approximate any function

Defining a Network Topology

- Decide the **network topology**: Specify # of units in the *input layer*, # of *hidden layers* (if > 1), # of units in *each hidden layer*, and # of units in the *output layer*
- Normalize the input values for each attribute measured in the training tuples to $[0.0—1.0]$
- One **input** unit per domain value, each initialized to 0
- **Output**, if for classification and more than two classes, one output unit per class is used
- Once a network has been trained and its accuracy is **unacceptable**, repeat the training process with a *different network topology* or a *different set of initial weights*

A Multi-Layer Feed-Forward Neural Network



$$w_j^{(k+1)} = w_j^{(k)} + \lambda(y_i - \hat{y}_i^{(k)})x_{ij}$$

Backpropagation

- Iteratively process a set of training tuples & compare the network's prediction with the actual known target value
- For each training tuple, the weights are modified to **minimize the mean squared error** between the network's prediction and the actual target value
- Modifications are made in the “**backwards**” direction: from the output layer, through each hidden layer down to the first hidden layer, hence “**backpropagation**”
- Steps
 - Initialize weights to small random numbers, associated with biases
 - Propagate the inputs forward (by applying activation function)
 - Backpropagate the error (by updating weights and biases)
 - Terminating condition (when error is very small, etc.)


Efficiency and Interpretability

- **Efficiency** of backpropagation: Each epoch (one iteration through the training set) takes $O(|D| * w)$, with $|D|$ tuples and w weights, but # of epochs can be exponential to n , the number of inputs, in worst case
- For easier comprehension: **Rule extraction** by network pruning
 - Simplify the network structure by removing weighted links that have the least effect on the trained network
 - Then perform link, unit, or activation value clustering
 - The set of input and activation values are studied to derive rules describing the relationship between the input and hidden unit layers
- **Sensitivity analysis**: assess the impact that a given input variable has on a network output. The knowledge gained from this analysis can be represented in rules

Neural Network as a Classifier


- Weakness
 - Long training time
 - Require a number of parameters typically best determined empirically, e.g., the network topology or “structure.”
 - Poor interpretability: Difficult to interpret the symbolic meaning behind the learned weights and of “hidden units” in the network
- Strength
 - High tolerance to noisy data
 - Ability to classify untrained patterns
 - Well-suited for continuous-valued inputs *and outputs*
 - Successful on an array of real-world data, e.g., hand-written letters
 - Algorithms are inherently parallel
 - Techniques have recently been developed for the extraction of rules from trained neural networks

Chapter 9. Classification: Advanced Methods

- Classification by Backpropagation
- Support Vector Machines 
- Additional Topics Regarding Classification
- Summary

Classification: A Mathematical Mapping

- **Classification:** predicts categorical class labels
 - E.g., Personal homepage classification
 - $x_i = (x_1, x_2, x_3, \dots)$, $y_i = +1$ or -1
 - x_1 : # of word “homepage”
 - x_2 : # of word “welcome”
- Mathematically, $x \in X = \mathbb{R}^n$, $y \in Y = \{+1, -1\}$,
 - We want to derive a function $f: X \rightarrow Y$
- Linear Classification
 - Binary Classification problem
 - Data above the red line belongs to class ‘x’
 - Data below red line belongs to class ‘o’
 - Examples: SVM, Perceptron, Probabilistic Classifiers



Discriminative Classifiers

- Advantages
 - Prediction accuracy is generally high
 - As compared to Bayesian methods – in general
 - Robust, works when training examples contain errors
 - Fast evaluation of the learned target function
 - Bayesian networks are normally slow
- Criticism
 - Long training time
 - Difficult to understand the learned function (weights)
 - Bayesian networks can be used easily for pattern discovery
 - Not easy to incorporate domain knowledge
 - Easy in the form of priors on the data or distributions


SVM—Support Vector Machines

- A relatively new classification method for both linear and nonlinear data
- It uses a nonlinear mapping to transform the original training data into a higher dimension
- With the new dimension, it searches for the linear optimal separating **hyperplane** (i.e., “decision boundary”)
- With an appropriate nonlinear mapping to a sufficiently high dimension, data from two classes can always be separated by a hyperplane
- SVM finds this hyperplane using **support vectors** (“essential” training tuples) and **margins** (defined by the support vectors)


SVM—History and Applications

- Vapnik and colleagues (1992)—groundwork from Vapnik & Chervonenkis' statistical learning theory in 1960s
- Features: training can be slow but accuracy is high owing to their ability to model complex nonlinear decision boundaries (margin maximization)
- Used for: classification and numeric prediction
- Applications:
 - handwritten digit recognition, object recognition, speaker identification, benchmarking time-series prediction tests


SVM—General Philosophy



SVM—Margins and Support Vectors



SVM—When Data Is Linearly Separable



Let data D be $(\mathbf{X}_1, y_1), \dots, (\mathbf{X}_{|D|}, y_{|D|})$, where \mathbf{X}_i is the set of training tuples associated with the class labels y_i

There are infinite lines (hyperplanes) separating the two classes but we want to find the best one (the one that minimizes classification error on unseen data)

*SVM searches for the hyperplane with the largest margin, i.e., **maximum marginal hyperplane** (MMH)*

SVM—Linearly Separable

- A separating hyperplane can be written as

$$\mathbf{W} \bullet \mathbf{X} + b = 0$$

where $\mathbf{W} = \{w_1, w_2, \dots, w_n\}$ is a weight vector and b a scalar (bias)

- For 2-D it can be written as

$$w_0 + w_1 x_1 + w_2 x_2 = 0$$

- The hyperplane defining the sides of the margin:

$$H_1: w_0 + w_1 x_1 + w_2 x_2 \geq 1 \quad \text{for } y_i = +1, \text{ and}$$

$$H_2: w_0 + w_1 x_1 + w_2 x_2 \leq -1 \quad \text{for } y_i = -1$$


- Any training tuples that fall on hyperplanes H_1 or H_2 (i.e., the sides defining the margin) are **support vectors**
- This becomes a **constrained (convex) quadratic optimization** problem: Quadratic objective function and linear constraints → *Quadratic Programming (QP)* → Lagrangian multipliers

Why Is SVM Effective on High Dimensional Data?

- The **complexity** of trained classifier is characterized by the # of support vectors rather than the dimensionality of the data
- The **support vectors** are the essential or critical training examples — they lie closest to the decision boundary (MMH)
- If all other training examples are removed and the training is repeated, the same separating hyperplane would be found
- The number of support vectors found can be used to compute an (upper) bound on the expected error rate of the SVM classifier, which is independent of the data dimensionality
- Thus, an SVM with a small number of support vectors can have good generalization, even when the dimensionality of the data is high

SVM—Linearly Inseparable

- Transform the original input data into a higher dimensional space



Example 6.8 Nonlinear transformation of original input data into a higher dimensional space. Consider the following example. A 3D input vector $\mathbf{X} = (x_1, x_2, x_3)$ is mapped into a 6D space Z using the mappings $\phi_1(\mathbf{X}) = x_1, \phi_2(\mathbf{X}) = x_2, \phi_3(\mathbf{X}) = x_3, \phi_4(\mathbf{X}) = (x_1)^2, \phi_5(\mathbf{X}) = x_1x_2$, and $\phi_6(\mathbf{X}) = x_1x_3$. A decision hyperplane in the new space is $d(\mathbf{Z}) = \mathbf{WZ} + b$, where \mathbf{W} and \mathbf{Z} are vectors. This is linear. We solve for \mathbf{W} and b and then substitute back so that we see that the linear decision hyperplane in the new (\mathbf{Z}) space corresponds to a nonlinear second order polynomial in the original 3-D input space,

$$\begin{aligned} d(\mathbf{Z}) &= w_1x_1 + w_2x_2 + w_3x_3 + w_4(x_1)^2 + w_5x_1x_2 + w_6x_1x_3 + b \\ &= w_1z_1 + w_2z_2 + w_3z_3 + w_4z_4 + w_5z_5 + w_6z_6 + b \end{aligned} \quad \blacksquare$$

- Search for a linear separating hyperplane in the new space

SVM: Different Kernel functions

- Instead of computing the dot product on the transformed data, it is math. equivalent to applying a kernel function $K(\mathbf{X}_i, \mathbf{X}_j)$ to the original data, i.e., $K(\mathbf{X}_i, \mathbf{X}_j) = \Phi(\mathbf{X}_i) \cdot \Phi(\mathbf{X}_j)$
- Typical Kernel Functions

Polynomial kernel of degree h : $K(\mathbf{X}_i, \mathbf{X}_j) = (\mathbf{X}_i \cdot \mathbf{X}_j + 1)^h$

Gaussian radial basis function kernel : $K(\mathbf{X}_i, \mathbf{X}_j) = e^{-\|\mathbf{X}_i - \mathbf{X}_j\|^2 / 2\sigma^2}$

Sigmoid kernel : $K(\mathbf{X}_i, \mathbf{X}_j) = \tanh(\kappa \mathbf{X}_i \cdot \mathbf{X}_j - \delta)$

- SVM can also be used for classifying multiple (> 2) classes and for regression analysis (with additional parameters)

SVM

- Consider the following data points. Please use SVM to train a classifier, and then classify these data points. Points with $a_i=1$ means this point is **support vector**. For example, point 1 (1,2) is the support vector, but point 5 (5,9) is not the support vector.
- Training data:

ID	a_i	x1	x2	y
1	1	1	2	1
2	1	2	1	-1
3	1	0	1	1
4	0	1	-2	-1
5	0	5	9	1
6	0	6	2	-1
7	0	3	9	1
8	0	7	1	-1

- Testing data:

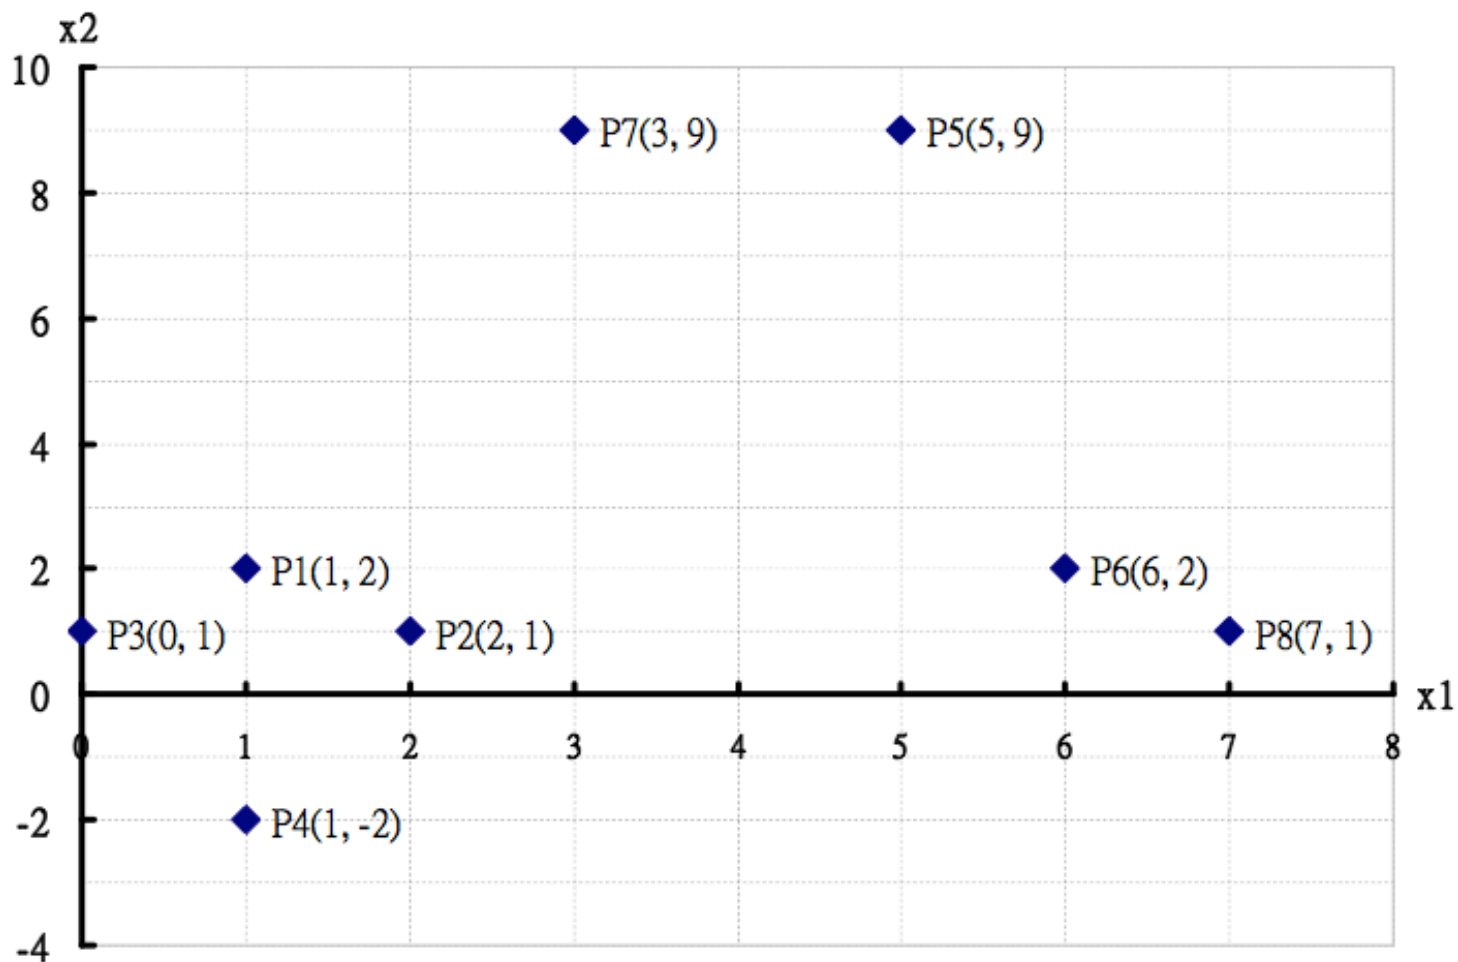
ID	x1	x2	y
9	2	5	
10	7	2	

SVM

- Question:
 - (a) Find the decision boundary, show detail calculation process.
 - (b) Use the decision boundary you found to classify the Testing data. Show all calculation process in detail, including the intermediate result and the formula you used.

SVM

- Answer:
- a) As the picture shows, P1, P2, P3 are support vectors.



SVM


- Suppose w is (w_1, w_2) . Since both $P1(1,2)$ and $P3(0,1)$ have $y = 1$, while $P2(2,1)$ has $y = -1$:
 - $w_1 * 1 + w_2 * 2 + b = 1$
 - $w_1 * 0 + w_2 * 1 + b = 1$
 - $w_1 * 2 + w_2 * 1 + b = -1$ $\Rightarrow w_1 = -1, w_2 = 1, b = 0$

then, the decision boundary is:

- $w_1 * x_1 + w_2 * x_2 + b = 0$
 $\Rightarrow -x_1 + x_2 = 0$

- Showed in the picture next page.

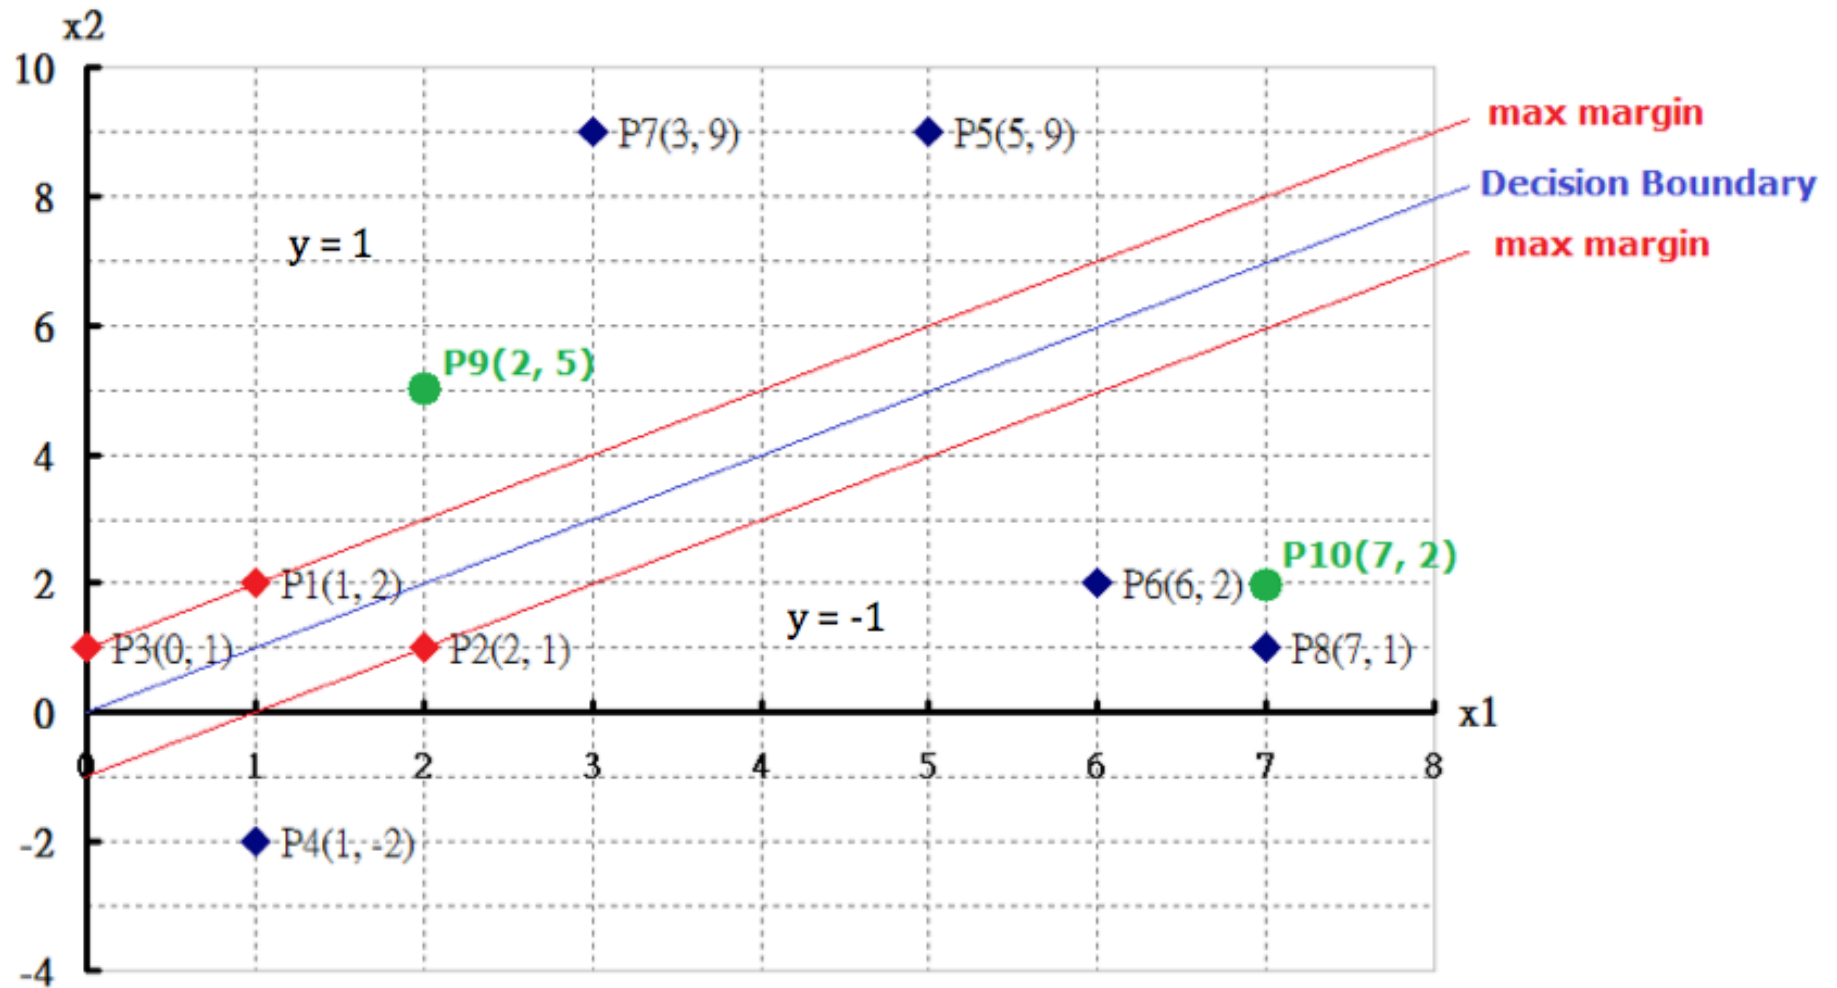
SVM



SVM

- b) Use the decision boundary to classify the testing data:
 - For the point P9 (2,5)
 $-x_1 + x_2 = -2 + 5 = 3 \geq 1$
So we choose $y = 1$
 - For the point P10 (7,2)
 $-x_1 + x_2 = -7 + 2 = -5 \leq -1$
So we choose $y = -1$
- Showed in the picture next page.

SVM



SVM vs. Neural Network

■ SVM

- Deterministic algorithm
- Nice generalization properties
- Hard to learn – learned in batch mode using quadratic programming techniques
- Using kernels can learn very complex functions


■ Neural Network

- Nondeterministic algorithm
- Generalizes well but doesn't have strong mathematical foundation
- Can easily be learned in incremental fashion
- To learn complex functions—use multilayer perceptron (nontrivial)

SVM Related Links

- SVM Website: <http://www.kernel-machines.org/>
- Representative implementations
 - **LIBSVM**: an efficient implementation of SVM, multi-class classifications, nu-SVM, one-class SVM, including also various interfaces with java, python, etc.
 - **SVM-light**: simpler but performance is not better than LIBSVM, support only binary classification and only in C
 - **SVM-torch**: another recent implementation also written in C

Chapter 9. Classification: Advanced Methods

- Classification by Backpropagation
- Support Vector Machines
- Additional Topics Regarding Classification 
- Summary

Multiclass Classification

- Classification involving more than two classes (i.e., > 2 Classes)
- Method 1. **One-vs.-all** (OVA): Learn a classifier one at a time
 - Given m classes, train m classifiers: one for each class
 - Classifier j : treat tuples in class j as *positive* & all others as *negative*
 - To classify a tuple \mathbf{X} , the set of classifiers vote as an ensemble
- Method 2. **All-vs.-all** (AVA): Learn a classifier for each pair of classes
 - Given m classes, construct $m(m-1)/2$ binary classifiers
 - A classifier is trained using tuples of the two classes
 - To classify a tuple \mathbf{X} , each classifier votes. \mathbf{X} is assigned to the class with maximal vote
- Comparison
 - All-vs.-all tends to be superior to one-vs.-all
 - Problem: Binary classifier is sensitive to errors, and errors affect vote count

Error-Correcting Codes for Multiclass Classification

- Originally designed to correct errors during data transmission for communication tasks by exploring data redundancy
- Example
 - A 7-bit codeword associated with classes 1-4
 - Given a unknown tuple \mathbf{X} , the 7-trained classifiers output: 0001010
 - Hamming distance: # of different bits between two codewords
 - $H(\mathbf{X}, C_1) = 5$, by checking # of bits between [1111111] & [0001010]
 - $H(\mathbf{X}, C_2) = 3$, $H(\mathbf{X}, C_3) = 3$, $H(\mathbf{X}, C_4) = 1$, thus C_4 as the label for \mathbf{X}
- Error-correcting codes can correct up to $(h-1)/h$ 1-bit error, where h is the minimum Hamming distance between any two codewords
- If we use 1-bit per class, it is equiv. to one-vs.-all approach, the code are insufficient to self-correct
- When selecting error-correcting codes, there should be good row-wise and col.-wise separation between the codewords


Class	Error-Corr. Codeword						
C_1	1	1	1	1	1	1	1
C_2	0	0	0	0	1	1	1
C_3	0	0	1	1	0	0	1
C_4	0	1	0	1	0	1	0

Semi-Supervised Classification

- Semi-supervised: Uses labeled and unlabeled data to build a classifier
- Self-training:
 - Build a classifier using the labeled data
 - Use it to label the unlabeled data, and those with the most confident label prediction are added to the set of labeled data
 - Repeat the above process
 - Adv: easy to understand; disadv: may reinforce errors
- Co-training: Use two or more classifiers to teach each other
 - Each learner uses a mutually independent set of features of each tuple to train a good classifier, say f_1
 - Then f_1 and f_2 are used to predict the class label for unlabeled data X
 - Teach each other: The tuple having the most confident prediction from f_1 is added to the set of labeled data for f_2 , & vice versa
- Other methods, e.g., joint probability distribution of features and labels

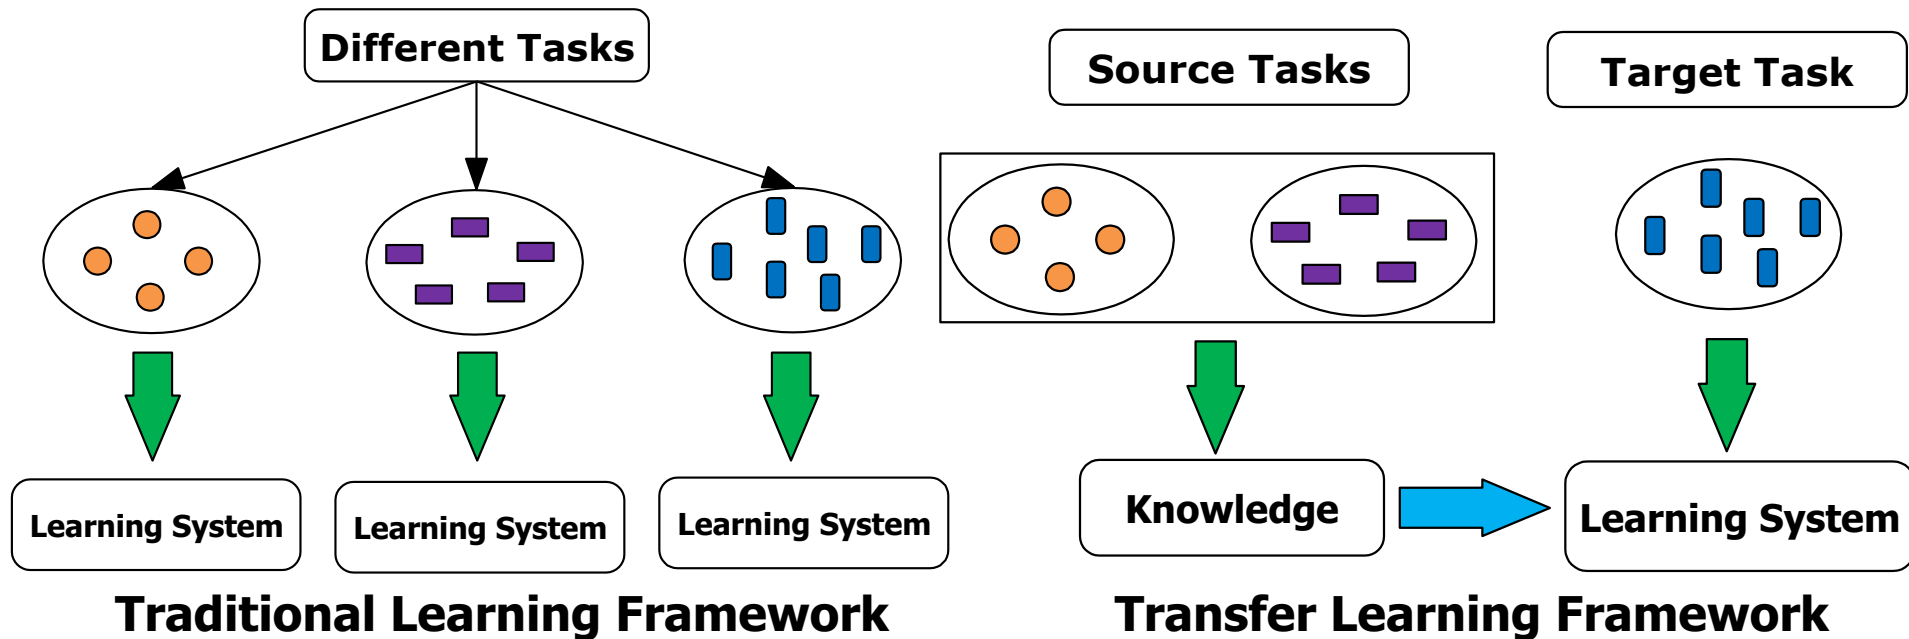
Active Learning

- Class labels are expensive to obtain
- Active learner: query human (oracle) for labels
- Pool-based approach: Uses a pool of unlabeled data
 - L: a small subset of D is labeled, U: a pool of unlabeled data in D
 - Use a query function to carefully select one or more tuples from U and request labels from an oracle (a human annotator)
 - The newly labeled samples are added to L, and learn a model
 - Goal: Achieve high accuracy using as few labeled data as possible
- Evaluated using *learning curves*: Accuracy as a function of the number of instances queried (# of tuples to be queried should be small)
- Research issue: How to choose the data tuples to be queried?
 - Uncertainty sampling: choose the least certain ones
 - Reduce *version space*, the subset of hypotheses consistent w. the training data
 - Reduce expected entropy over U: Find the greatest reduction in the total number of incorrect predictions



Transfer Learning: Conceptual Framework


- Transfer learning: Extract knowledge from one or more source tasks and apply the knowledge to a target task
- Traditional learning: Build a new classifier for each new task
- Transfer learning: Build new classifier by applying existing knowledge learned from source tasks



Transfer Learning: Methods and Applications

- Applications: Especially useful when data is outdated or distribution changes, e.g., Web document classification, e-mail spam filtering
- *Instance-based transfer learning*: Reweight some of the data from source tasks and use it to learn the target task
- TrAdaBoost (Transfer AdaBoost)
 - Assume source and target data each described by the same set of attributes (features) & class labels, but rather diff. distributions
 - Require only labeling a small amount of target data
 - Use source data in training: When a source tuple is misclassified, reduce the weight of such tuples so that they will have less effect on the subsequent classifier
- Research issues
 - Negative transfer: When it performs worse than no transfer at all
 - Heterogeneous transfer learning: Transfer knowledge from different feature space or multiple source domains
 - Large-scale transfer learning

Chapter 9. Classification: Advanced Methods

- Classification by Backpropagation
- Support Vector Machines
- Additional Topics Regarding Classification
- Summary 

Summary

- Effective and advanced classification methods
 - Bayesian belief network (probabilistic networks)
 - Backpropagation (Neural networks)
 - Support Vector Machine (SVM)
 - Pattern-based classification
 - Other classification methods: lazy learners (KNN, case-based reasoning), genetic algorithms, rough set and fuzzy set approaches
- Additional Topics on Classification
 - Multiclass classification
 - Semi-supervised classification
 - Active learning
 - Transfer learning

SVM—Introductory Literature

- “Statistical Learning Theory” by Vapnik: extremely hard to understand, containing many errors too.
- C. J. C. Burges. [A Tutorial on Support Vector Machines for Pattern Recognition](#). *Knowledge Discovery and Data Mining*, 2(2), 1998.
 - Better than the Vapnik’s book, but still written too hard for introduction, and the examples are so not-intuitive
- The book “An Introduction to Support Vector Machines” by N. Cristianini and J. Shawe-Taylor
 - Also written hard for introduction, but the explanation about the mercer’s theorem is better than above literatures
- The neural network book by Haykins
 - Contains one nice chapter of SVM introduction


Notes about SVM— Introductory Literature

- “Statistical Learning Theory” by **Vapnik**: difficult to understand, containing many errors.
- C. J. C. **Burges**. A Tutorial on Support Vector Machines for Pattern Recognition. *Knowledge Discovery and Data Mining*, 2(2), 1998.
 - Easier than Vapnik’s book, but still not introductory level; the examples are not so intuitive
- The book An Introduction to Support Vector Machines by **Cristianini and Shawe-Taylor**
 - Not introductory level, but the explanation about Mercer’s Theorem is better than above literatures
- Neural Networks and Learning Machines by **Haykin**
 - Contains a nice chapter on SVM introduction


Associative Classification Can Achieve High Accuracy and Efficiency (Cong et al. SIGMOD05)

Dataset	RCBT	CBA	IRG Classifier	C4.5 family			SVM
				single tree	bagging	boosting	
AML/ALL (ALL)	91.18%	91.18%	64.71%	91.18%	91.18%	91.18%	97.06%
Lung Cancer(LC)	97.99%	81.88%	89.93%	81.88%	96.64%	81.88%	96.64%
Ovarian Cancer(OC)	97.67%	93.02%	-	97.67%	97.67%	97.67%	97.67%
Prostate Cancer(PC)	97.06%	82.35%	88.24%	26.47%	26.47%	26.47%	79.41%
Average Accuracy	95.98%	87.11%	80.96%	74.3%	77.99%	74.3%	92.70%


Table 2: Classification Results



(a) ALL-AML leukemia



(b) Lung Cancer



(c) Ovarian Cancer

A Closer Look at CMAR

- **CMAR** (Classification based on Multiple Association Rules: Li, Han, Pei, ICDM'01)
- Efficiency: Uses an enhanced FP-tree that maintains the distribution of class labels among tuples satisfying each frequent itemset
- Rule pruning whenever a rule is inserted into the tree
 - Given two rules, R_1 and R_2 , if the antecedent of R_1 is more general than that of R_2 and $\text{conf}(R_1) \geq \text{conf}(R_2)$, then prune R_2
 - Prunes rules for which the rule antecedent and class are not positively correlated, based on a χ^2 test of statistical significance
- Classification based on generated/pruned rules
 - If only *one rule* satisfies tuple X , assign the class label of the rule
 - If a *rule set* S satisfies X , CMAR
 - divides S into groups according to class labels
 - uses a weighted χ^2 measure to find the strongest group of rules, based on the statistical correlation of rules within a group
 - assigns X the class label of the strongest group